

Introduction to Statistical Inquiry

Adel Mohammadpour
University of Calgary
adel.mohammadpour@ucalgary.ca

SeyedParsa Hosseinipour Rafsanjani
parsahosseini2001@gmail.com

April 25, 2025

This book is published under a Creative Commons BY-SA license (CC BY-SA) version 4.0. This means that this book can be reused, remixed, retained, revised and redistributed (including commercially) as long as appropriate credit is given to the authors. If you remix, or modify the original version of this open textbook, you must redistribute all versions of this open textbook under the same license - CC BY-SA.

<https://creativecommons.org/licenses/by-sa/4.0/>

Table of Contents

1	Probability and Risk	1
1.1	Introduction	1
1.2	Random Experiment and Sample Space	2
1.3	Event	3
1.4	Set Theory Operations on Events	4
1.5	What is Probability?	7
2	Probability and Accuracy	13
2.1	Probability	13
2.2	Conditional Probability & Bayes' Theorem	13
2.3	Prevalence, Relative Risk, Sensitivity, and Specificity	13
3	Descriptive Statistics	15
3.1	Introduction to Data	15
3.2	Exploratory Data Analysis	15
4	Probability Distributions	17
4.1	Random Variables	17
4.2	Binomial Distribution	17
4.3	Normal Distribution	17
4.4	Sampling Distribution and Central Limit Theorem	17
5	Estimation	19
5.1	Point Estimation	19
5.2	Confidence Interval	19
6	Estimation	21
6.1	Hypothesis Testing for One Sample	21
6.2	Hypothesis Testing for Two Variances	21
6.3	Hypothesis for Two Means	21
6.4	Hypothesis for Two Proportions	21
7	Hypothesis Testing	23
8		25
9	Probability and Accuracy	27
10		29

11	31
12	33
13	35
14	37
15	39
16	41
17 Probability and Accuracy	43
17.1 Probability	43
17.2 Conditional Probability & Bayes' Theorem	43
17.3 Prevalence, Relative Risk, Sensitivity, and Specificity	43
18 Descriptive Statistics	45
18.1 Introduction to Data	45
18.2 Exploratory Data Analysis	45
19 Probability Distributions	47
19.1 Random Variables	47
19.2 Binomial Distribution	47
19.3 Normal Distribution	47
19.4 Sampling Distribution and Central Limit Theorem	47

1. Probability and Risk

“When there are but two players, your theory which proceeds by combinations is very just. But when there are three, I believe I have a proof that it is unjust that you should proceed in any other manner than the one I have.”

– Pascal’s letter to Fermat¹

1.1

Introduction

While scientists have always tried to understand the universe with technology and explain it with complete certainty, this is not always possible. In other words, there is no other way but to accept chance as part of our lives.

The concept of chance has been extensively explored across elementary, professional, and philosophical literature, serving as a key motivation for our study. While chance often appears unpredictable and devoid of structure, mathematicians have long sought to define it through rules and systematic frameworks. Not surprisingly, gamblers were among the first to seek systematic frameworks for understanding their games - probing the mechanics of luck, wins and losses. In 1654, a Parisian gambler named Antoine Gombaud (alias Chevalier de Méré), posed critical questions about winning probabilities to two of the era’s greatest mathematicians: Blaise Pascal and Pierre de Fermat. Through their correspondence, the two mathematicians initiated the development of modern probability theory. However, Gerolamo Cardano and Galileo Galilei, two Italian scholars, had also made significant contributions that captured the interest of Italian gamblers.

After years of development, the Russian mathematician Andrey Kolmogorov introduced the standard probability axioms in 1933, establishing the rigorous foundations of modern probability theory. Today, probability theory serves as a fundamental tool across diverse fields including social sciences,

¹from <https://www.york.ac.uk/depts/maths/histstat/pascal.pdf>

medicine, biology, machine learning, physics, and countless other applications. The following sections provide rigorous definitions of key concepts needed to establish a precise theoretical foundation.

1.2

Random Experiment and Sample Space

Suppose we want to conduct an experiment for which we know all possible outcomes. Assuming the experimental conditions remain constant each time, if each realization of this experiment produces exactly one outcome, we call it a **random experiment**. We denote the set of all possible outcomes of such an experiment by S , called the **sample space**, which is a nonempty set.

Example 1.2.1 *Consider the experiment of tossing a coin where the experimental conditions are controlled such that the coin lands on heads or tails, with no other possible outcomes. For instance the surface is chosen so that it won't land on edge. This is a random experiment in which we are interested in observing whether the coin lands heads or tails when viewed from above after it comes to rest. Hence, denoting the outcome of observing heads by H and tails by T , the sample space of this experiment is $S = \{H, T\}$.*

Example 1.2.2 *Consider an experiment where a ball is drawn from an urn containing one blue, one green, and one red ball. The experimenter cannot see inside the urn when making each draw. This is a random experiment in which we are interested in the color of the drawn ball. Hence, denoting the outcome of observing blue, green and red ball by B, G and R respectively, the sample space of this experiment is $S = \{B, G, R\}$.*

Example 1.2.3 *Consider the experiment of drawing a card from a deck. Each time a card is drawn, it is returned to the deck and the deck is shuffled, so the conditions of each experiment remains the same. If we are interested in the suit of the drawn cards, the sample space is $S = \{\clubsuit, \diamondsuit, \spadesuit, \heartsuit\}$. But if we are interested in the rank of the cards, the sample space becomes $S = \{A, 1, 2, 3, 4, 5, 6, 7, 8, 9, J, Q, K\}$.*

In some cases, as the following example demonstrates, the sample space may be infinite:

Example 1.2.4 *A coin is repeatedly tossed under the experimental conditions of Example 1.2.1 until heads appears. There are infinitely many possible outcomes: If the first coin toss results in heads, the experiment is immediately terminated. If it's tails, the coin is tossed again. The experiment terminates if heads appears for the second time. If not, the experiment continues until heads is observed. So the sample space in this case is $S = \{H, TH, TTH, TTTH, \dots\}$.*

The first outcome is the case where heads appears on the first coin toss. The second outcome corresponds to heads appearing on the second coin toss, and so on.

Event

Each outcome of a random experiment is an element of its sample space, S . In doing these experiments, we are interested in observing some specific outcomes, or a subset of S . This subset is called an **event** and is denoted by E . But we know that in each realization of a random experiment, only one element $e \in S$ is observed. If $e \in E$, we say E has occurred and if $e \in S - E$, we say it has not occurred.

Example 1.3.1 *In the experiment of throwing a fair die, $S = \{1, 2, 3, 4, 5, 6\}$. In case the experimenter is interested in the event of observing an even number, the event of interest would be $E = \{2, 4, 6\}$, a subset of S . If instead they are interested in the event of observing an odd number, the event of interest would become $E = \{1, 3, 5\}$, again a subset of S .*

Note that in the previous example, we omitted certain details about the experimental conditions and our observations of interest. Henceforth, unless explicitly stated otherwise, we make the following standard assumptions:

- Each realization of a random experiment is performed under identical conditions.
- For a die roll, the random experiment of interest is observing the uppermost face after landing on the ground.
- For a coin toss, the random experiment of interest is whether it lands heads (H) or tails (T).

If $E = S$, the event is called a **sure event**, since each element of S is in E and so observing any outcome means the event has occurred. On the other hand, if $E = \emptyset$, the event is called an **impossible event**, since it contains no element at all.

Example 1.3.2 *In Example 1.3.1, observing a number greater than 6 is an impossible event, while observing a number less than 7 is a sure event.*

An event consisting of only a single element of the sample space S is called an **elementary event** or an **atomic event**.

Example 1.3.3 *In Example 1.3.1, $\{1\}$, $\{2\}$, $\{3\}$, $\{4\}$, $\{5\}$ and $\{6\}$ are all elementary events.*

For two events E and F , if every element of E is also in F , we say that E is a **subevent** of F , or in set notation, $E \subset F$.

Example 1.3.4 *In Example 1.3.1, if E is the event "observing a number greater than 5" and F is the event "observing a number greater than 4", then $E = \{6\}$ and $F = \{5, 6\}$, and thus $E \subset F$. It is trivial that the occurrence of E implies the occurrence of F , but not conversely. For instance, if the die roll outcome is 5, F occurs but E does not.*

Two events are called **equal events** if they consist of same elements. By definition, if $E \subset F$ and $F \subset E$, then E and F are equal events and we write $E = F$.

Example 1.3.5 In Example 1.2.3, suppose the sample space is $S = \{A, 1, 2, 3, 4, 5, 6, 7, 8, 9, J, Q, K\}$. If E is the event "observing a number less than 5 in Clubs" and F is the event "observing a number less than 5 in Diamonds", then $E = \{1, 2, 3, 4\}$ and $F = \{1, 2, 3, 4\}$, and thus $E = F$. Trivially, the occurrence of E implies the occurrence of F , and vice versa.

1.4

Set Theory Operations on Events

We denoted the set of all outcomes of a random experiment by S , and we saw that we may be interested in a subset of this set, $E \subset S$. Now consider two events $E, F \subset S$.

1.4.1 Union of Two Events

The **union** of two events E and F consists of all elements that are in E or in F and is denoted by $E \cup F$. The occurrence of E or F results in the occurrence of $E \cup F$.

Example 1.4.1 In Example 1.3.1, we denote the event of observing an even number by E and the event of observing a prime number by F . So $E = \{2, 4, 6\}$ and $F = \{2, 3, 5\}$. The union of these two events consists of all the elements in E or F , meaning $E \cup F = \{2, 3, 4, 5, 6\}$. So if the outcome of a die roll is 4, we say the event "observing a number which is even or prime" occurs, since 4 is an even number.

If it's 5, we also say that this event occurs, since 5 is a prime number.

Now suppose we observe 2 in a die roll. We again say that $E \cup F$ occurs. The important thing to note here is that this "or" is an inclusive or, which means $E \cup F$ is actually "observing a number that is even or prime or both". In this case, 2 is both an even and a prime number.

1.4.2 Intersection of Two Events

The **intersection** of two events E and F consists of all elements in both E and F and is denoted by $E \cap F$. In order for $E \cap F$ to occur, both E and F has to occur.

Example 1.4.2 Consider the events E and F in Example 1.4.1. $E \cap F$ is the event "observing a number that is both even and prime". In other words, $E \cap F$ consists of all elements that are in both E and F , and so are both even and prime. Hence, $E \cap F = \{2\}$.

So if 3 is observed in rolling a die, $E \cap F$ does not occur since 3 is not even while it is a prime number. If 4 is observed, $E \cap F$ again does not occur, since 4 is even but not prime. The only acceptable observation to say $E \cap F$ occurs is 2, since only 2 is both even and prime among all the elements of E and F .

If $E \cap F = \emptyset$, we say the two events are **disjoint** or **mutually exclusive**. For instance, if we denote the event "observing an odd number" by O , $E \cap O = \emptyset$, meaning E and O are incompatible. In other words, a single die roll cannot simultaneously result in both an even and an odd number.

Not that events A_1, A_2, \dots are disjoint (or mutually exclusive) if $A_i \cap A_j = \emptyset$ for $i \neq j$.

1.4.3 Difference of Two Events

The **difference** $E - F$ is the event containing all elements in E but not in F .

Example 1.4.3 In Example 1.4.1, $E - F$ is the event "observing a number that is even but not prime". Thus $E - F = \{4, 6\}$. $3 \notin E - F$ since it is prime, and $2 \notin E - F$ because while it is an even number, it is also a prime number. So observing 3 or 2 upon rolling a die means $E - F$ does not occur. For $E - F$ to occur, we have to observe either 4 or 6.

1.4.4 Complement of an Event

The **complement** of an event E is the difference $S - E$, containing all elements in sample space that are not in E . The complement of E is denoted by E^c , E' or \bar{E} .

Example 1.4.4 In Example 1.4.1, E^c is the event "observing a number that is not even", which means $E^c = \{1, 3, 5\}$. So E^c occurs only when the observed number on the die is odd.

Note that "observing a number that is not even" does not mean any number that is not even, but rather any number in the sample space $S = \{1, 2, 3, 4, 5, 6\}$ that is not even.

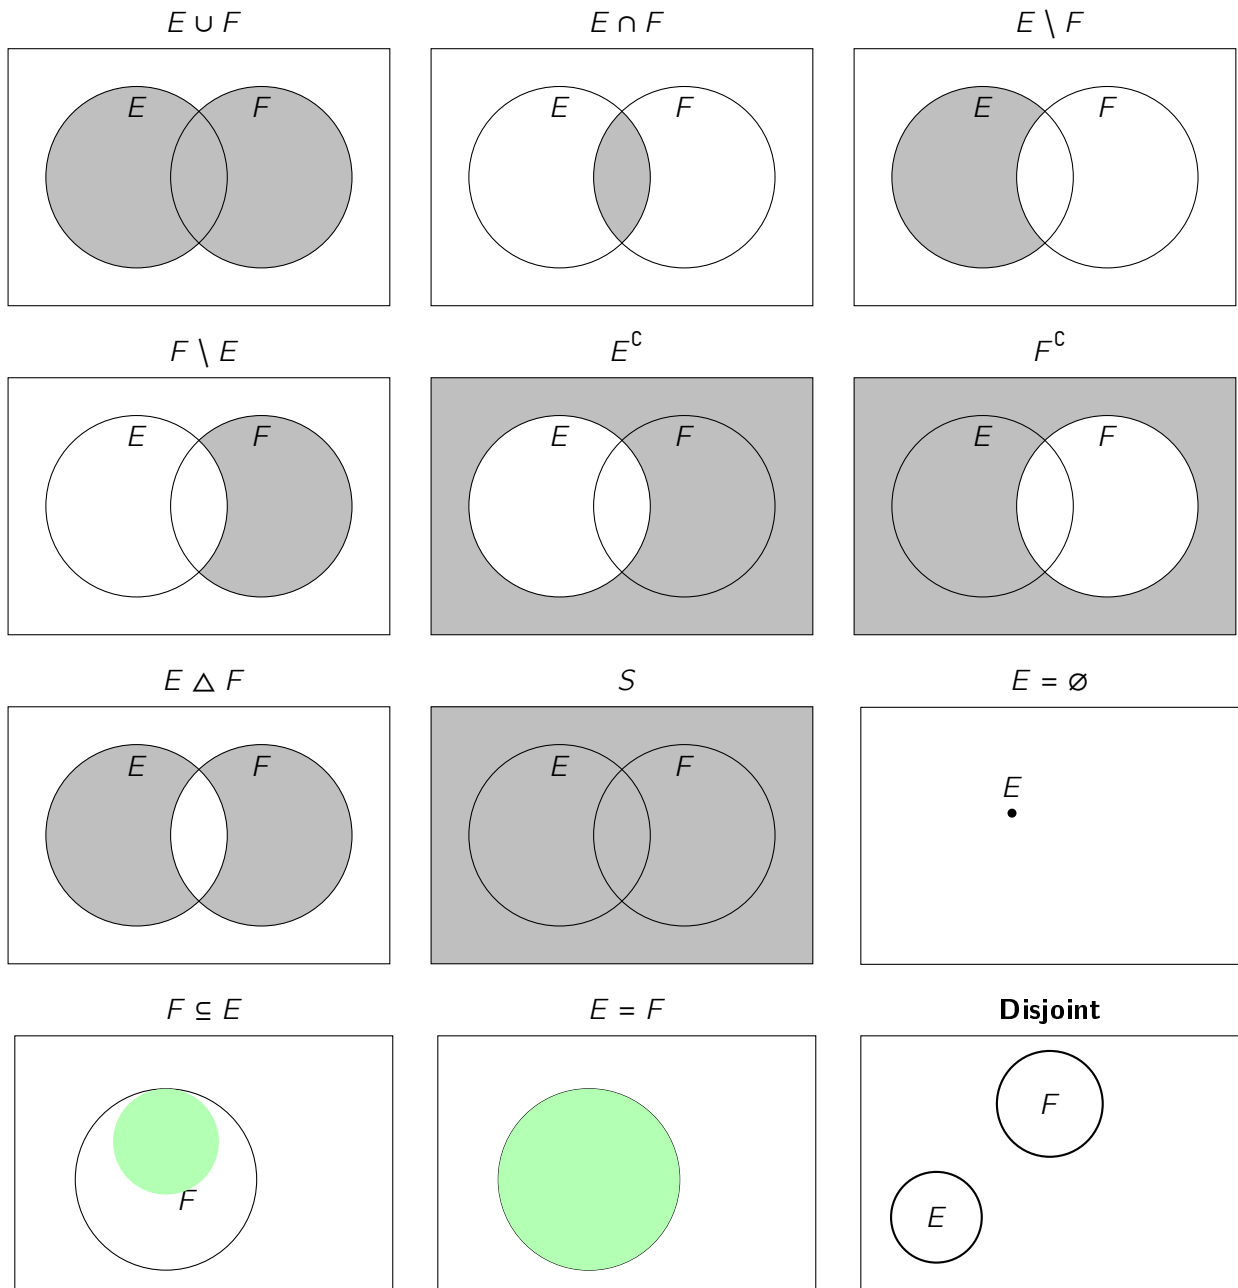
1.4.5 Symmetric Difference of Two Events

The **symmetric difference** of two events E and F consists of all elements in E or F but not in both and is denoted by $E \Delta F$. By definition, $E \Delta F = F \Delta E$ and this is why this operation is symmetric.

Example 1.4.5 In Example 1.4.1, $E \Delta F = \{3, 4, 5, 6\}$, containing numbers in the sample space that are even or prime but not both. Since 2 is both even and prime, it follows that $2 \in E \cup F$ but $2 \notin E \Delta F$.

Note the exclusive "or" in this operation, unlike the inclusive case discussed in Example 1.4.1.

The set operations and key concepts for two events E and F are visually represented on this page. In these diagrams, the sample space is depicted as a rectangle, events E and F as two circles, and the operation or concept specified above each rectangle is highlighted through coloring.



What is Probability?

There are three prominent ways to define probability. In the three following subsections, these three definitions are explored. Then an axiomatic definition of probability is given which is consistent with these three definitions.

1.5.1 Classical Definition of Probability

In this definition, each elementary event of the sample space is considered to be equally likely to occur. For example in rolling a fair die, all the events $\{1\}$, $\{2\}$, $\{3\}$, $\{4\}$, $\{5\}$ and $\{6\}$ have the same probability. Or in tossing a coin (the random experiment in Example 1.2.1), the elementary events $\{H\}$ and $\{T\}$ are all equally likely to happen, which is consistent with our intuition of 50% chance of occurring for each event.

This interpretation of Probability goes back to Pierre-Simon Laplace, who wrote:

The probability of an event is the ratio of the number of cases favorable to it, to the number of all cases possible when nothing leads us to expect that any one of these cases should occur more than any other, which renders them, for us, equally possible. ²

While we can assign a probability to each elementary event when the sample space consists of finitely many elements, it is not possible when the elementary events are not equally likely to occur or the sample space is infinite like we saw in Example 1.2.4.

1.5.2 Frequentist Definition of Probability

Consider a random experiment in which we are interested in the occurrence of event E . Suppose this experiment is repeated n times under identical experimental conditions, and event E occurs r times in total. r and $\frac{r}{n}$ are said to be the **frequency** and **relative frequency** of E in these n trials, respectively.

As the number of trials n increases, the frequency r and consequently the relative frequency $\frac{r}{n}$ changes as well. However, empirical observations show that the relative frequency converges to a constant value, which in this interpretation is defined as the probability of E .

The frequentist interpretation of probability may trace its earliest conceptual origins to Aristotle, who wrote

the probable is that which for the most part happens ³

²Laplace, Théorie analytique des probabilités, retrieved from https://en.wikipedia.org/wiki/Classical_definition_of_probability

³Aristotle, Rhetoric, retrieved from https://en.wikipedia.org/wiki/Frequentist_probability

The frequentist interpretation of probability rose to dominance during the 19th century, becoming the foundation of classic statistical inference.

Let us reconsider the random experiment described in Example 1.2.2. According to the classical interpretation of probability, the probability of drawing a blue ball is $\frac{1}{3}$. However, if we don't know the urn's contents, we can determine the probability of drawing a red ball through a simple experiment. By repeatedly drawing balls from the urn and calculating the relative frequency of drawing blue balls, we observe that this ratio converges to $\frac{1}{3}$ as the number of trials increases, which is consistent with the classical definition of probability. But the frequentist interpretation offers another advantage: Suppose there are two blue balls, one green ball and one red ball in the urn. Here, while the sample space remains $\{B, G, R\}$, the elementary events now have different probabilities because there is an additional blue ball in the urn! So while we cannot calculate the probability of drawing a blue ball using the classical definition, we can determine it through the frequentist approach by performing many trials and observing the convergence of relative frequency.

While this interpretation sounds like a better approach to defining probability, it faces several limitations. Consider, for instance, estimating the probability of precipitation occurring tomorrow. There is, in fact, only one tomorrow; we cannot conduct multiple trials of "tomorrow" to count rainy occurrences and determine their relative frequency. Moreover, maintaining truly identical experimental conditions is practically impossible in most real-world scenarios. Another fundamental challenge lies in determining how many trials are sufficient for the relative frequency to converge to a stable probability value.

1.5.3 Epistemic Definition of Probability

In this interpretation of probability, each observer assigns a subjective probability to an event based on their prior beliefs. Reconsider the urn example from the previous subsection. The experimenter draws a ball from the urn without revealing it to the observers.

An observer who saw the experimenter add an extra blue ball to the urn assigns a probability of $\frac{2}{4}$ to the ball being blue. Another observer, who previously knew the urn contained one blue, one green, and one red ball, assigns a probability of $\frac{1}{3}$. The experimenter, however, knows with certainty whether the ball is blue or not.

In the frequentist approach, the ball is either blue or not, and the probability is determined through repeated trials. Thus, frequentists don't assign probabilities to single drawn balls. In the epistemic perspective, however, each individual assigns a probability based on their existing knowledge.

1.5.4 Axiomatic Definition of Probability

Just like Euclidean geometry, where theorems are derived from a set of axioms taken as true⁴, modern probability theory is built upon axioms proposed by the Russian mathematician Andrey Kolmogorov in

⁴<https://www.math.brown.edu/tbanchof/Beyond3d/chapter9/section01.html>

1933.

A **probability measure** on a sample space S is a function $P(\cdot)$ that assigns to each event $E \subset S$ a real number $P(A)$, called the probability of A , satisfying the following three **axioms of probability**:

- for each event $E \subset S$, the probability of A is non-negative, meaning $P(E) \geq 0$.
- the probability of S is 1, meaning $P(S) = 1$.
- if E_1, E_2, \dots is a sequence of disjoint events, meaning $E_i \cap E_j = \emptyset$ for each $i \neq j$, then $P(E_1 \cup E_2 \cup \dots) = P(E_1) + P(E_2) + \dots$.

The first axiom states that probabilities cannot be negative, which aligns with our intuition. The second axiom indicates that since every experimental outcome belongs to the sample space, the event S must occur with probability 1. The third axiom establishes that for any collection of disjoint events, the probability of their union equals the sum of their individual probabilities.

Different interpretations of probability remain consistent with the axiomatic definition. Consider the frequentist approach:

- Relative frequencies are non-negative, satisfying the first axiom.
- The relative frequency of S equals 1 in any number of trials, since every outcome belongs to the sample space.
- For mutually exclusive events, the relative frequency of their union equals the sum of their individual relative frequencies, as each outcome is counted only once.

Next, we prove some theorems using the axioms of probability.

Theorem 1.5.1 *The probability of impossible event, \emptyset , is zero.*

Proof: Consider the sequence of events E_1, E_2, \dots where $E_1 = S$ and $E_i = \emptyset$ for $i > 1$. Since the events in this sequence are mutually exclusive and $S = \bigcup_{i=1}^{\infty} E_i$, by the third axiom we have

$$P(S) = P\left(\bigcup_{i=1}^{\infty} A_i\right) = \sum_{i=1}^{\infty} P(E_i) = P(S) + \sum_{i=2}^{\infty} P(\emptyset) = P(S) + P\left(\bigcup_{i=2}^{\infty} \emptyset\right) = P(S) + P(\emptyset)$$

Cancelling $P(S)$ from both sides we obtain $P(\emptyset) = 0$. □

Theorem 1.5.2 *For finitely many disjoint events E_1, E_2, \dots, E_n , we have*

$$P\left(\bigcup_{i=1}^n E_i\right) = \sum_{i=1}^n P(E_i)$$

Proof: Consider the sequence of events E_1, E_2, \dots where $E_i = \emptyset$ for $i > n$. Since all events in this sequence are mutually exclusive and $\bigcup_{i=1}^n E_i = \bigcup_{i=1}^{\infty} E_i$, by applying the third axiom and [Theorem 1.5.1](#) we have

$$\begin{aligned}
 P\left(\bigcup_{i=1}^n E_i\right) &= P\left(\bigcup_{i=1}^{\infty} E_i\right) = \sum_{i=1}^{\infty} P(E_i) \\
 &= \sum_{i=1}^n P(E_i) + \sum_{i=n+1}^{\infty} P(E_i) \\
 &= \sum_{i=1}^n P(E_i) + P\left(\bigcup_{i=n+1}^{\infty} E_i\right) \\
 &= \sum_{i=1}^n P(E_i) + P(\emptyset) \\
 &= \sum_{i=1}^n P(E_i)
 \end{aligned}$$

□

Theorem 1.5.3 $P(E^c) = 1 - P(E)$.

Proof: E and E^c are mutually exclusive and $S = E \cup E^c$ since every element of the sample space is either in E or not. Thus, using the second axiom and [Theorem 1.5.2](#), we have

$$\begin{aligned}
 1 &= P(S) = P(E \cup E^c) = P(E) + P(E^c) \\
 P(E^c) &= 1 - P(E)
 \end{aligned}$$

□

Theorem 1.5.4 For any two events $F, E \subset S$, the equality $P(F - E) = P(F) - P(E \cap F)$ holds.

Proof: It can be shown that $F = (F - E) \cup (E \cap F)$ where $F - E$ and $E \cap F$ are mutually exclusive. By [Theorem 1.5.2](#):

$$\begin{aligned}
 P(F) &= P(F - E) + P(E \cap F) \\
 P(F - E) &= P(F) - P(E \cap F)
 \end{aligned}$$

□

Corollary 1.5.4.1 If $E \subset F$, then $E \cap F = E$ and thus $P(F - E) = P(F) - P(E)$.

Corollary 1.5.4.2 *If $E \subset F$, the first axiom and Corollary 1.5.4.1 yield $P(E) \leq P(F)$. Replacing F with S and applying the second axiom gives $P(E) \leq 1$. Combined with the first axiom, this proves that for any event E ,*

$$0 \leq P(E) \leq 1$$

Theorem 1.5.5 *For any two events $E, F \subset S$, the equality $P(E \cup F) = P(E) + P(F) - P(F \cap E)$ holds.*

Proof: It can be shown that $E \cap F = E \cup (F - E)$ where E and $F - E$ are mutually exclusive. By Theorem 1.5.2 and Theorem 1.5.4:

$$P(E \cup F) = P(E) + P(F - E) = P(E) + P(F) - P(E \cap F)$$

□

Glossary

2. Probability and Accuracy

2.1

Probability

2.2

Conditional Probability & Bayes' Theorem

2.3

Prevalence, Relative Risk, Sensitivity, and Specificity

Glossary

3. Descriptive Statistics

3.1

Introduction to Data

3.2

Exploratory Data Analysis

Glossary

4. Probability Distributions

4.1

Random Variables

4.2

Binomial Distribution

4.3

Normal Distribution

4.4

Sampling Distribution and Central Limit Theorem

Glossary

5. Estimation

5.1

Point Estimation

5.2

Confidence Interval

Glossary

6. Estimation

6.1 _____

Hypothesis Testing for One Sample

6.2 _____

Hypothesis Testing for Two Variances

6.3 _____

Hypothesis for Two Means

6.4 _____

Hypothesis for Two Proportions

Glossary

7. Hypothesis Testing

Glossary

- **anecdotal evidence:** Evidence, often personal, that is collected casually rather than by a well-designed study.
- **population:** A group we are interested in studying. “Population” often refers to a group of people, but the term is used for other subjects, too.
- **cross-sectional study:** A study that collects data about a population at a particular point in time.
- **cycle:** In a repeated cross-sectional study, each repetition of the study is called a cycle.
- **longitudinal study:** A study that follows a population over time, collecting data from the same group repeatedly.
- **record:** In a dataset, a collection of information about a single person or other subject.
- **respondent:** A person who responds to a survey.
- **sample:** The subset of a population used to collect data.
- **representative:** A sample is representative if every member of the population has the same chance of being in the sample.
- **oversampling:** The technique of increasing the representation of a sub-population in order to avoid errors due to small sample sizes.
- **raw data:** Values collected and recorded with little or no checking, calculation or interpretation.
- **recode:** A value that is generated by calculation and other logic applied to raw data.
- **data cleaning:** Processes that include validating data, identifying errors, translating between data types and representations, etc.

8.

Glossary

- **linear fit:** a line intended to model the relationship between variables.
- **least squares fit:** A model of a dataset that minimizes the sum of squares of the residuals.
- **residual:** The deviation of an actual value from a model.
- **goodness of fit:** A measure of how well a model fits data.
- **coefficient of determination:** A statistic intended to quantify goodness of fit.
- **sampling weight:** A value associated with an observation in a sample that indicates what part of the population it represents.

9. Probability and Accuracy

Glossary

- **regression**: One of several related processes for estimating parameters that fit a model to data.
- **dependent variables**: The variables in a regression model we would like to predict. Also known as endogenous variables.
- **explanatory variables**: The variables used to predict or explain the dependent variables. Also known as independent, or exogenous, variables.
- **simple regression**: A regression with only one dependent and one explanatory variable.
- **multiple regression**: A regression with multiple explanatory variables, but only one dependent variable.
- **linear regression**: A regression based on a linear model.
- **ordinary least squares**: A linear regression that estimates parameters by minimizing the squared error of the residuals.
- **spurious relationship**: A relationship between two variables that is caused by a statistical artifact or a factor, not included in the model, that is related to both variables.
- **control variable**: A variable included in a regression to eliminate or “control for” a spurious relationship.
- **proxy variable**: A variable that contributes information to a regression model indirectly because of a relationship with another factor, so it acts as a proxy for that factor.
- **categorical variable**: A variable that can have one of a discrete set of unordered values.
- **join**: An operation that combines data from two DataFrames using a key to match up rows in the two frames.

- **data mining**: An approach to finding relationships between variables by testing a large number of models.
- **logistic regression**: A form of regression used when the dependent variable is boolean.
- **Poisson regression**: A form of regression used when the dependent variable is a non-negative integer, usually a count.
- **odds**: An alternative way of representing a probability, p , as the ratio of the probability and its complement, $p/(1 - p)$.

10.

Glossary

- **time series:** A dataset where each value is associated with a timestamp, often a series of measurements and the times they were collected.
- **window:** A sequence of consecutive values in a time series, often used to compute a moving average.
- **moving average:** One of several statistics intended to estimate the underlying trend in a time series by computing averages (of some kind) for a series of overlapping windows.
- **rolling mean:** A moving average based on the mean value in each window.
- **exponentially-weighted moving average (EWMA):** A moving average based on a weighted mean that gives the highest weight to the most recent values, and exponentially decreasing weights to earlier values.
- **span:** A parameter of EWMA that determines how quickly the weights decrease.
- **serial correlation:** Correlation between a time series and a shifted or lagged version of itself.
- **lag:** The size of the shift in a serial correlation or autocorrelation.
- **autocorrelation:** A more general term for a serial correlation with any amount of lag.
- **autocorrelation function:** A function that maps from lag to serial correlation.
- **stationary:** A model is stationary if the parameters and the distribution of residuals does not change over time.

11.

Glossary

- **survival analysis:** A set of methods for describing and predicting lifetimes, or more generally time until an event occurs.
- **survival curve:** A function that maps from a time, t , to the probability of surviving past t .
- **hazard function:** A function that maps from t to the fraction of people alive until t who die at t .
- **Kaplan-Meier estimation:** An algorithm for estimating hazard and survival functions.
- **cohort:** a group of subjects defined by an event, like date of birth, in a particular interval of time.
- **cohort effect:** a difference between cohorts.
- **NBUE:** A property of expected remaining lifetime, “New better than used in expectation.”
- **UBNE:** A property of expected remaining lifetime, “Used better than new in expectation.”

12.

Glossary

- **distribution**: The values that appear in a sample and the frequency of each.
- **histogram**: A mapping from values to frequencies, or a graph that shows this mapping.
- **frequency**: The number of times a value appears in a sample.
- **mode**: The most frequent value in a sample, or one of the most frequent values.
- **normal distribution**: An idealization of a bell-shaped distribution; also known as a Gaussian distribution.
- **uniform distribution**: A distribution in which all values have the same frequency.
- **tail**: The part of a distribution at the high and low extremes.
- **central tendency**: A characteristic of a sample or population; intuitively, it is an average or typical value.
- **outlier**: A value far from the central tendency.
- **spread**: A measure of how spread out the values in a distribution are.
- **summary statistic**: A statistic that quantifies some aspect of a distribution, like central tendency or spread.
- **variance**: A summary statistic often used to quantify spread.
- **standard deviation**: The square root of variance, also used as a measure of spread.
- **effect size**: A summary statistic intended to quantify the size of an effect like a difference between groups.
- **clinically significant**: A result, like a difference between groups, that is relevant in practice.

13.

Glossary

- **Probability mass function (PMF)**: a representation of a distribution as a function that maps from values to probabilities.
- **probability**: A frequency expressed as a fraction of the sample size.
- **normalization**: The process of dividing a frequency by a sample size to get a probability.
- **index**: In a pandas DataFrame, the index is a special column that contains the row labels.

14.

Glossary

- **percentile rank:** The percentage of values in a distribution that are less than or equal to a given value.
- **percentile:** The value associated with a given percentile rank.
- **cumulative distribution function (CDF):** A function that maps from values to their cumulative probabilities. $CDF(x)$ is the fraction of the sample less than or equal to x .
- **inverse CDF:** A function that maps from a cumulative probability, p , to the corresponding value.
- **median:** The 50th percentile, often used as a measure of central tendency.
- **interquartile range:** The difference between the 75th and 25th percentiles, used as a measure of spread.
- **quantile:** A sequence of values that correspond to equally spaced percentile ranks; for example, the quartiles of a distribution are the 25th, 50th and 75th percentiles.
- **replacement:** A property of a sampling process. “With replacement” means that the same value can be chosen more than once; “without replacement” means that once a value is chosen, it is removed from the population.

15.

Glossary

- **empirical distribution:** The distribution of values in a sample.
- **analytic distribution:** A distribution whose CDF is an analytic function.
- **model:** A useful simplification. Analytic distributions are often good models of more complex empirical distributions.
- **interarrival time:** The elapsed time between two events.
- **complementary CDF:** A function that maps from a value, x , to the fraction of values that exceed x , which is $1 - CDF(x)$.
- **standard normal distribution:** The normal distribution with mean 0 and standard deviation 1.
- **normal probability plot:** A plot of the values in a sample versus random values from a standard normal distribution.

Glossary

- **Probability density function (PDF):** The derivative of a continuous CDF, a function that maps a value to its probability density.
- **Probability density:** A quantity that can be integrated over a range of values to yield a probability. If the values are in units of cm, for example, probability density is in units of probability per cm.
- **Kernel density estimation (KDE):** An algorithm that estimates a PDF based on a sample.
- **discretize:** To approximate a continuous function or distribution with a discrete function. The opposite of smoothing.
- **raw moment:** A statistic based on the sum of data raised to a power.
- **central moment:** A statistic based on deviation from the mean, raised to a power.
- **standardized moment:** A ratio of moments that has no units.
- **skewness:** A measure of how asymmetric a distribution is.
- **sample skewness:** A moment-based statistic intended to quantify the skewness of a distribution.
- **Pearson's median skewness coefficient:** A statistic intended to quantify the skewness of a distribution based on the median, mean, and standard deviation.
- **robust:** A statistic is robust if it is relatively immune to the effect of outliers.

17. Probability and Accuracy

17.1

Probability

17.2

Conditional Probability & Bayes' Theorem

17.3

Prevalence, Relative Risk, Sensitivity, and Specificity

Glossary

- **scatter plot:** A visualization of the relationship between two variables, showing one point for each row of data.
- **jitter:** Random noise added to data for purposes of visualization.
- **saturation:** Loss of information when multiple points are plotted on top of each other.
- **correlation:** A statistic that measures the strength of the relationship between two variables.
- **standardize:** To transform a set of values so that their mean is 0 and their variance is 1.
- **standard score:** A value that has been standardized so that it is expressed in standard deviations from the mean.
- **covariance:** A measure of the tendency of two variables to vary together.

- **rank:** The index where an element appears in a sorted list.
- **randomized controlled trial:** An experimental design in which subjects are divided into groups at random, and different groups are given different treatments.
- **treatment group:** A group in a controlled trial that receives some kind of intervention.
- **control group:** A group in a controlled trial that receives no treatment, or a treatment whose effect is known.
- **natural experiment:** An experimental design that takes advantage of a natural division of subjects into groups in ways that are at least approximately random.

18. Descriptive Statistics

18.1

Introduction to Data

18.2

Exploratory Data Analysis

Glossary

- **estimation**: The process of inferring the parameters of a distribution from a sample.
- **estimator**: A statistic used to estimate a parameter.
- **mean squared error (MSE)**: A measure of estimation error.
- **root mean squared error (RMSE)**: The square root of MSE, a more meaningful representation of typical error magnitude.
- **maximum likelihood estimator (MLE)**: An estimator that computes the point estimate most likely to be correct.
- **bias (of an estimator)**: The tendency of an estimator to be above or below the actual value of the parameter, when averaged over repeated experiments.
- **sampling error**: Error in an estimate due to the limited size of the sample and variation due to chance.
- **sampling bias**: Error in an estimate due to a sampling process that is not representative of the population.

- **measurement error:** Error in an estimate due to inaccuracy collecting or recording data.
- **sampling distribution:** The distribution of a statistic if an experiment is repeated many times.
- **standard error:** The RMSE of an estimate, which quantifies variability due to sampling error (but not other sources of error).
- **confidence interval:** An interval that represents the expected range of an estimator if an experiment is repeated many times.

19. Probability Distributions

19.1

Random Variables

19.2

Binomial Distribution

19.3

Normal Distribution

19.4

Sampling Distribution and Central Limit Theorem

Glossary

- **hypothesis testing:** The process of determining whether an apparent effect is statistically significant.
- **test statistic:** A statistic used to quantify an effect size.
- **null hypothesis:** A model of a system based on the assumption that an apparent effect is due to chance.
- **p-value:** The probability that an effect could occur by chance.

- **statistically significant:** An effect is statistically significant if it is unlikely to occur by chance.
- **permutation test:** A way to compute p-values by generating permutations of an observed dataset.
- **resampling test:** A way to compute p-values by generating samples, with replacement, from an observed dataset.
- **two-sided test:** A test that asks, "What is the chance of an effect as big as the observed effect, positive or negative?"
- **one-sided test:** A test that asks, "What is the chance of an effect as big as the observed effect, and with the same sign?"
- **chi-squared test:** A test that uses the chi-squared statistic as the test statistic.
- **false positive:** The conclusion that an effect is real when it is not.
- **false negative:** The conclusion that an effect is due to chance when it is not.
- **power:** The probability of a positive test if the null hypothesis is false.

